



Syncany

Secure open-source file sync

Contact the Syncany Team
[via team@syncany.org](mailto:team@syncany.org)

Agenda



- **What is Syncany?**

- How does it work?

- Focus Topics:
 - Deduplication & Versioning
 - Cryptography Concept
 - Continuous Integration & Packaging

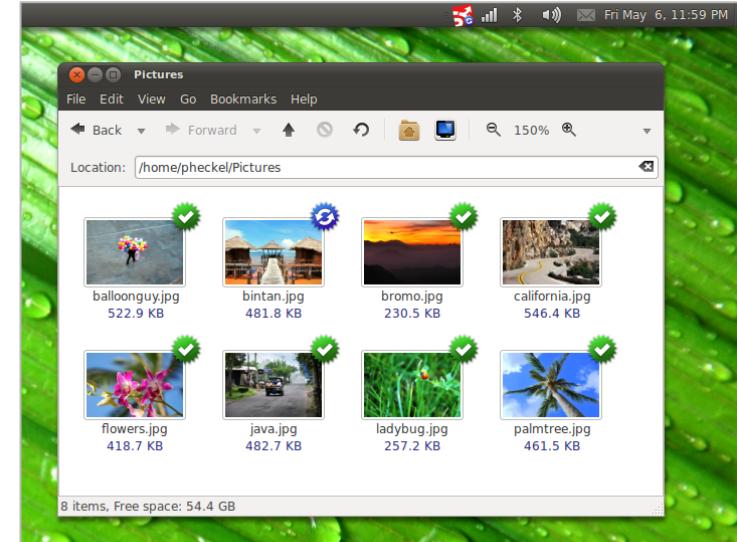
What is Syncany?



Syncany is an open source Dropbox-like file sync and backup tool. It synchronizes files and folders between computers either manually or automatically. Users can define certain folders on their machine and keep them in sync with friends or colleagues.

So what makes Syncany **awesome**? What makes it **different**?

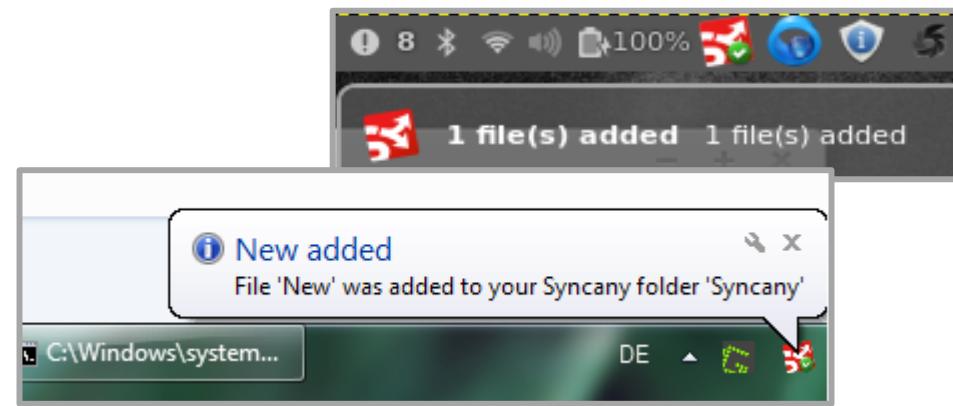
- Syncany is **open-source** software
- **Use-your-own** storage (FTP, S3, WebDAV, ...)
- Files are **encrypted** before upload (don't trust anybody but yourself)
- Files are **deduplicated** locally (space savings on remote storage)
- Files are intelligently **versioned** (restoring old versions is easy)



Manual version control-like sync

```
/bin/bash
pheckel@platop ~/Syncany $ sy status
? Germany.jpg
? The Netherlands.jpg
pheckel@platop ~/Syncany $ sy up
A Germany.jpg
A The Netherlands.jpg
Sync up finished.
pheckel@platop ~/Syncany $ sy down
A Canada.jpg
Sync down finished.
pheckel@platop ~/Syncany $
```

Automatic Dropbox-like sync





Storage Plugins

> Use-your-own storage, or even better: use any storage!

Storage plugins:

- Amazon S3
- Dropbox
- Flickr
- FTP
- RAID0
- Local (mounted drives; e.g. via FUSE)
- OpenStack Swift
- SFTP
- Samba / CIFS
- PHP / REST (3rd party, incubating)
- WebDAV

*(Easy to develop your own
in about 200-300 lines of code)*

Other plugins:

- Graphical User Interface
- Simple Web Interface (PoC, looking for dev)

The image shows two windows from the Syncany application. The top window is titled 'Preferences' and has a 'Plugins' tab selected. It lists several installed plugins: Dropbox (0.4.3-alpha), Flickr (0.4.1-alpha, User), and FTP (0.4.0-alpha, User). There are also dropdown menus for 'Global' and '0.4.0-alpha' with download icons. The bottom window is titled 'Syncany' and contains 'WebDAV settings'. It includes fields for 'URL' (https://dav.syncany.org/test/repo1), 'Username' (repo1), and 'Password (not displayed)' (redacted). Navigation buttons at the bottom include '< Previous', 'Next >', and 'Cancel'.

Project History



First attempt: ~ What I learned: Never advertise a prototype!

- Started in 2011, announcement of PoC on WebUpd8 (popular blog) and Ubuntu Podcast in May 2011
- 15,000 visitors in announcement week, 400 developers in mailing list, translated in 30 languages by excited helpers, >2,200 followers on Twitter; then big drop in interest due to non-functioning software.

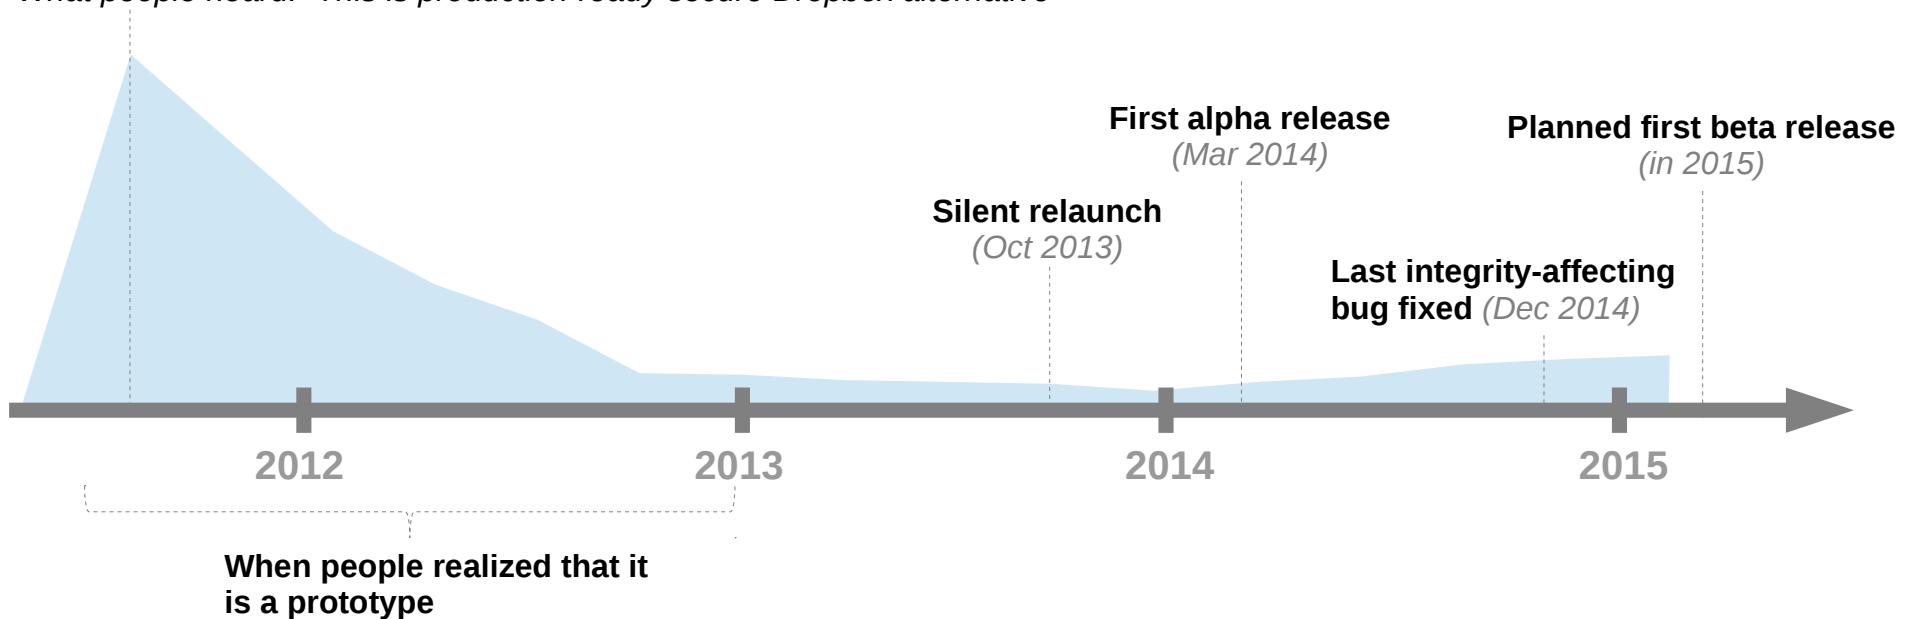
Second attempt: ~ Doing it right: Make a stable core, then build on top of that!

- Silent relaunch in October 2013 (no new announcement yet)
- 3-4 active developers, ~30 contributors, IRC channel with 25-30 people

When we announced the project (May 2011)

What we said: "This is a proof-of-concept, we need people to help"

What people heard: "This is production-ready secure Dropbox alternative"



Agenda



- What is Syncany?

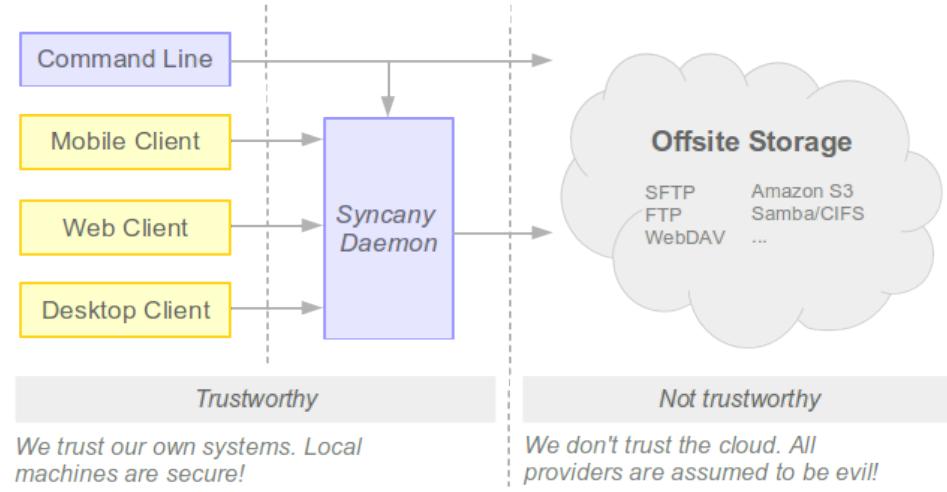
- **How does it work?**

- Focus Topics:
 - Deduplication & Versioning
 - Cryptography Concept
 - Continuous Integration & Packaging

How does it work?



High Level Architecture



Components

- **Core library:** Behavior logic (e.g. up/down operation), persistence, plugin architeture
- **Command line interface (CLI):** `$ sy <up|down|...>`
- **Background daemon:** REST/WS API for GUI/Web/Mobile interface
- **Plugins:** Storage plugins & user interface plugins

Core Operations

UpOperation ~ part of core library



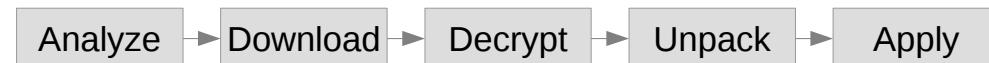
CLI

`$ sy up`

Daemon/GUI

When local changes occur & every 2min

DownOperation ~ part of core library



`$ sy down`

When remote changes occur & every 2 min

Agenda



- What is Syncany?
- How does it work?
- Focus Topics:
 - **Deduplication & Versioning**
 - Cryptography Concept
 - Continuous Integration & Packaging

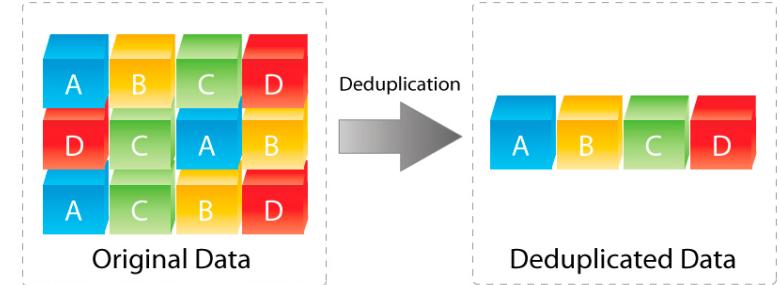


Deduplication & Versioning

> Syncany uses source-side, in-line, fixed-size deduplication

Data deduplication identifies duplicate sequences of bytes (chunks) and only stores **one copy of that sequence**.

If a chunk appears again in the same file (or in another file), deduplication only **stores a reference to this chunk** instead of its actual contents.



Distinction to compression

Not bound to location, date or time of files; stateful approach: local database/index necessary

Characteristics

- Location: Target-side vs. **Source-side**
- Time: **In-line** vs. Post-process
- Chunking method: whole file vs. **fixed-size** vs. variable-size vs. format-aware chunking

Highlighted:
Concepts in Syncany

Chunking methods

Methods have different impact on performance, I/O, CPU and RAM usage

- Whole file chunking: Only deduplicate entire files, files are not split into chunks.
- Fixed-size chunking: Files are split into chunks at fixed offsets, e.g. at 32 KB, 64 KB, 96 KB, ...
- Variable-size chunking: Files are split into chunks using “*sliding window*”-based chunking method (such as TTTD) and a *fingerprinting* method to determine the break points, e.g. using rolling checksum algorithms such as Adler-32 or Rabin.
- ...

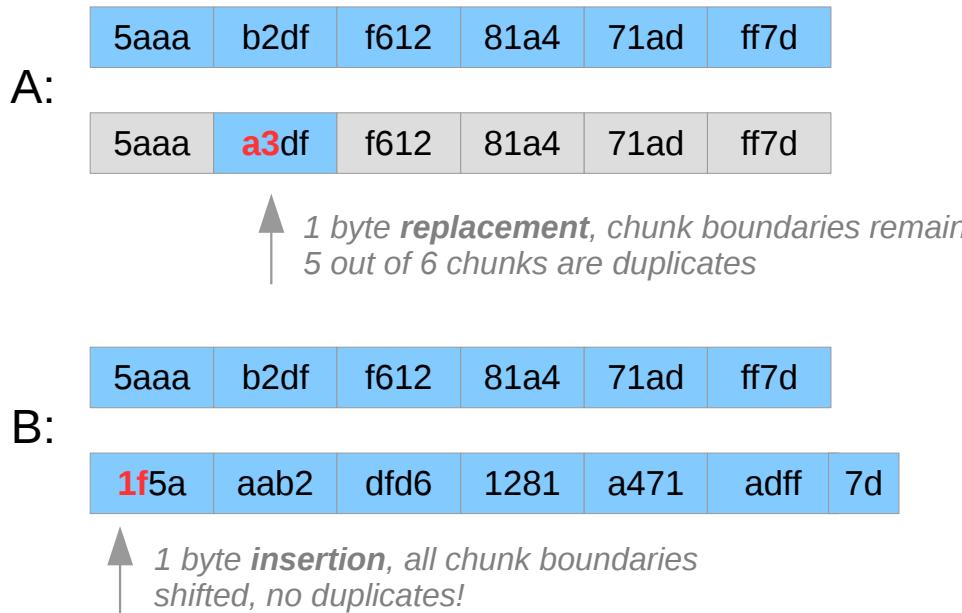
Details on
next slide

Deduplication & Versioning

> Syncany uses fixed-size chunking, but should use variable-size chunking



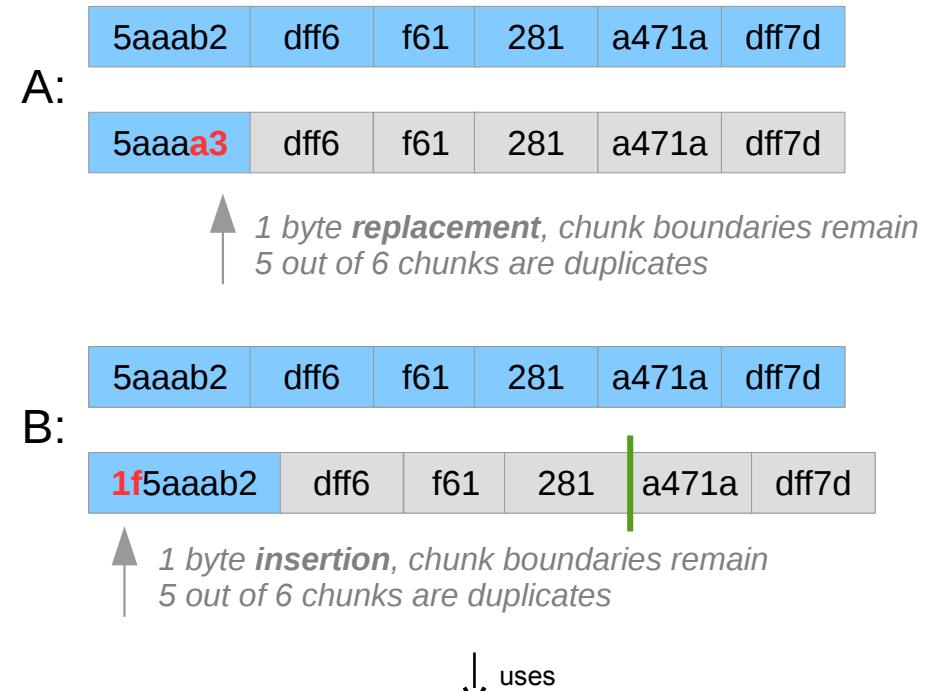
Fixed-Size Chunking



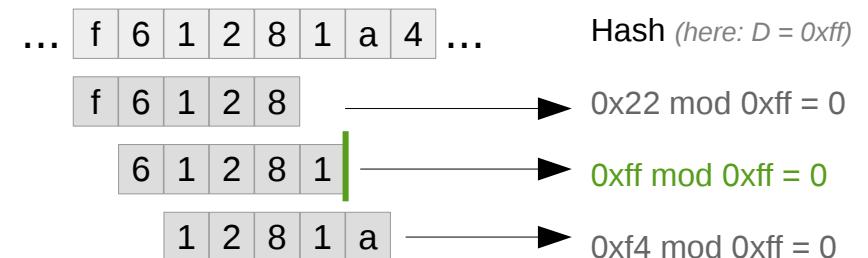
Legend:	
5aaa	New chunk
5aaa	Duplicate chunk

Chunk break condition:
 $f \bmod D = 0$
f is hash value
D is divisor

Variable-Size Chunking



Sliding window based rolling hash



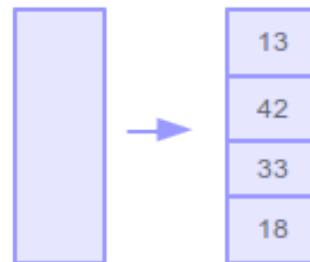


Deduplication & Versioning

> Syncany: Fixed-size, client-side dedup. with additional multichunking

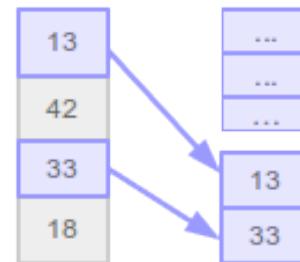
1. File chunking

Index all files and divide them into unique chunks



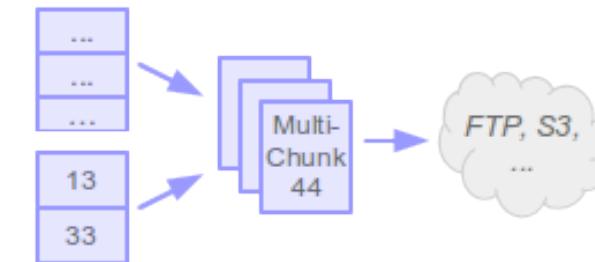
2. Create multichunks

Compare chunks to database and combine new chunks to multichunks



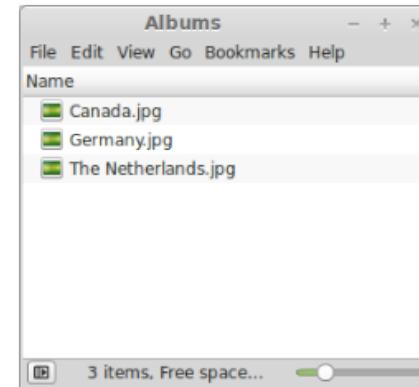
3. Upload multichunks

Encrypt and upload multichunks to offsite storage using a storage plugin



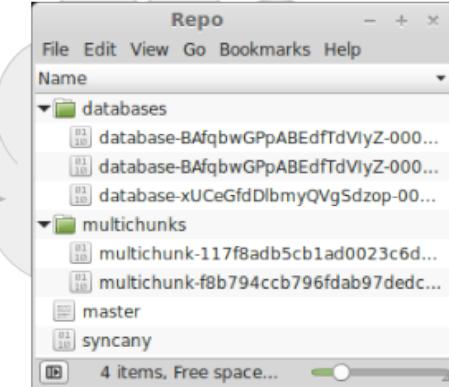
Current implementation:

- Fixed-size chunking, 512 KB chunks
- 4 MB multichunks



Local Syncany Folder

Syncany's local files are not encrypted, but are encrypted before upload.



Offsite Repository

Syncany stores file data in multichunks, and metadata in database files



Deduplication & Versioning

> Dedup. makes versioning easy, Syncany just keeps track of chunks

File 1a8a8dfeea4

			Chunks			
 Presentation.ppt (version 1)	NEW		5aa.. b2d.. f61.. 81a.. 71a.. ff7..		+ 3.1 MB	
 Presentation.ppt (version 2)	CHANGED		5aa.. 88a.. f61.. 81a.. 71a.. ff7..		+ 512 KB	
 Syncany.ppt (version 3)	MOVED		5aa.. 88a.. f61.. 81a.. 71a.. ff7..		+/- 0 Bytes	
 Syncany.ppt (version 4)	DELETED				+/- 0 Bytes	

File dfa818e192ad

 Products.ppt (version 1)	NEW	Chunks	5aa.. b2d.. 9ae.. e78.. aa1.. 79e..	+ 2.1 MB
--	------------	--------	-------------------------------------	----------



*Free copies (copies of files add 0 byte to the repo)
Small file changes add minimal additional repo space
Restoring old/deleted files is easy*

Agenda



- What is Syncany?
- How does it work?
- Focus Topics:
 - Deduplication & Versioning
 - **Cryptography Concept**
 - Continuous Integration & Packaging



Cryptography Concept

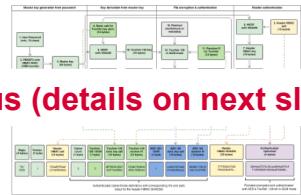
> Assumptions and Algorithms

Security Assumptions

- Storage provider and network infrastructure cannot be trusted
- The local machine is secure
- Users sharing a repository trust each other completely

Security Concepts

- Pre-upload encryption ← Focus (details on next slide!)
- HTTPS-only daemon/server



Algorithms:

- **PBKDF2 with HMAC-SHA1, 1mm rounds, 512-bit salt**
Password-Based Key Derivation Function 2, part of PKCS #5 v2.0
Applies SHA1 multiple times on the password and salt to derive a key from the password
Used in Syncany: per-repository master key generation

```
SecretKeyFactory factory = SecretKeyFactory.getInstance("PBKDF2WithHmacSHA1");
KeySpec pbeKeySpec = new PBEKeySpec(password.toCharArray(), salt, 1000000, 512);
SecretKey masterKey = factory.generateSecret(pbeKeySpec);
```

- **HKDF with SHA256, 96-bit salt**
HMAC-based Extract-and-Expand Key Derivation Function (HKDF), RFC5869

Creates a derived key from the master key (input key material) and a salt
Used in Syncany: per-file key is derived from the master key

```
HKDFBytesGenerator hkdf = new HKDFBytesGenerator(new SHA256Digest());
hkdf.init(new HKDFParameters(inputKeyMaterial, inputSalt, "Syncany..."));
byte[] derivedKey = new byte[outputKeySize / 8];
hkdf.generateBytes(derivedKey, 0, derivedKey.length);
```

- **HMAC-SHA256, key derived with HKDF, 96-bit salt**

Keyed-Hash Message Authentication Code (HMAC), RFC2104
Calculates a MAC using a cryptographic hash function (here: SHA256)
Used in Syncany: per-file header authentication

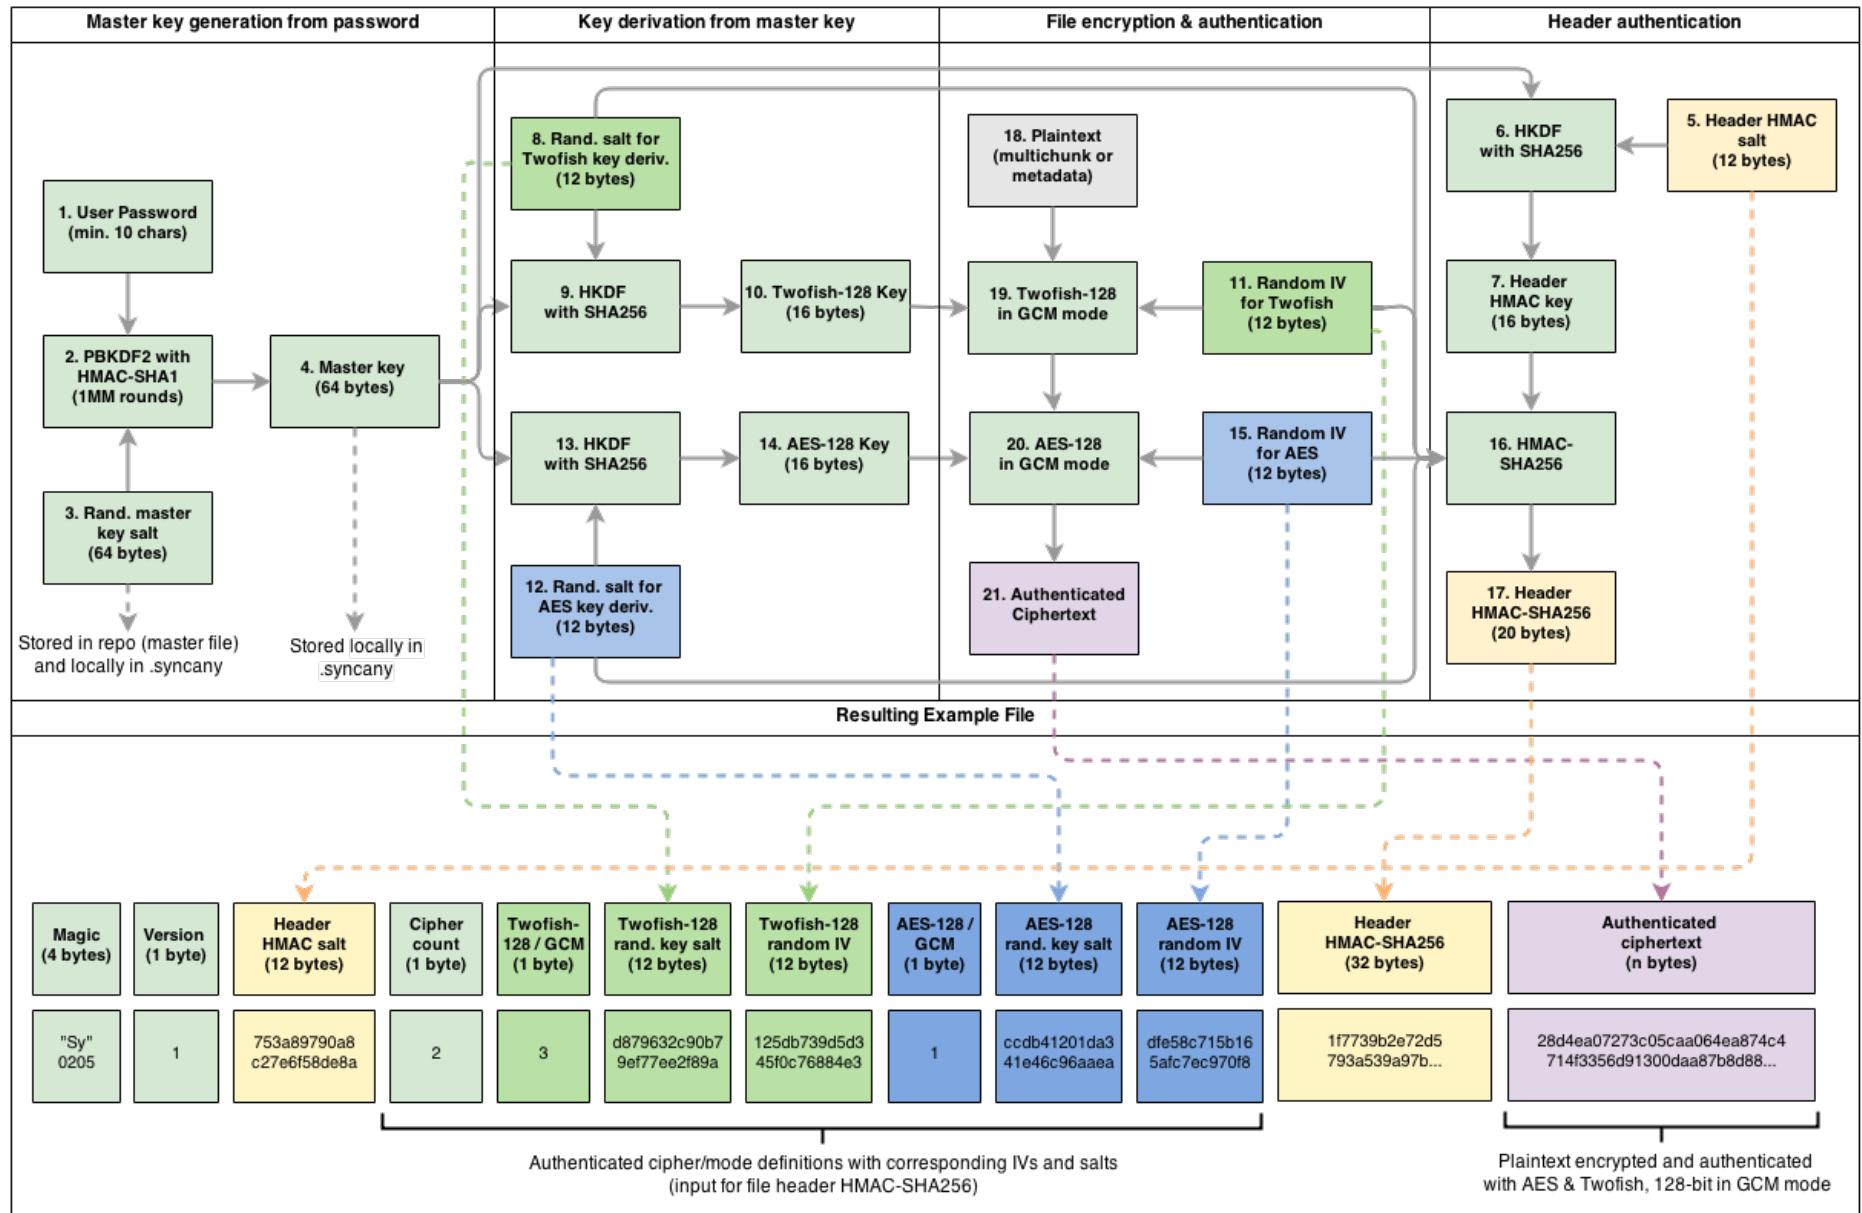
```
headerHmac = Mac.getInstance("HmacSHA256", "BC");
headerHmac.init(hmacSecretKey); // via HKDF
headerHmac.update(headerBytes);
outputStream.write(headerHmac.doFinal());
```

- **Twofish/AES 128-bit in GCM mode, key derived with HKDF, 96-bit salt**

Symmetric encryption scheme, used in the Galois Counter Mode (GCM), an Authenticated Encryption with Associated Data (AEAD)
Used in Syncany: per-file encryption; AES and Twofish in nested

Cryptography Concept

> Encryption file format and used algorithms



Agenda



- What is Syncany?
- How does it work?
- Focus Topics:
 - Deduplication & Versioning
 - Cryptography Concept
 - **Continuous Integration & Packaging**



Continuous Integration & Packaging

> Commit to `develop` to create a snapshot, or to `master` for a release

Workflow components:

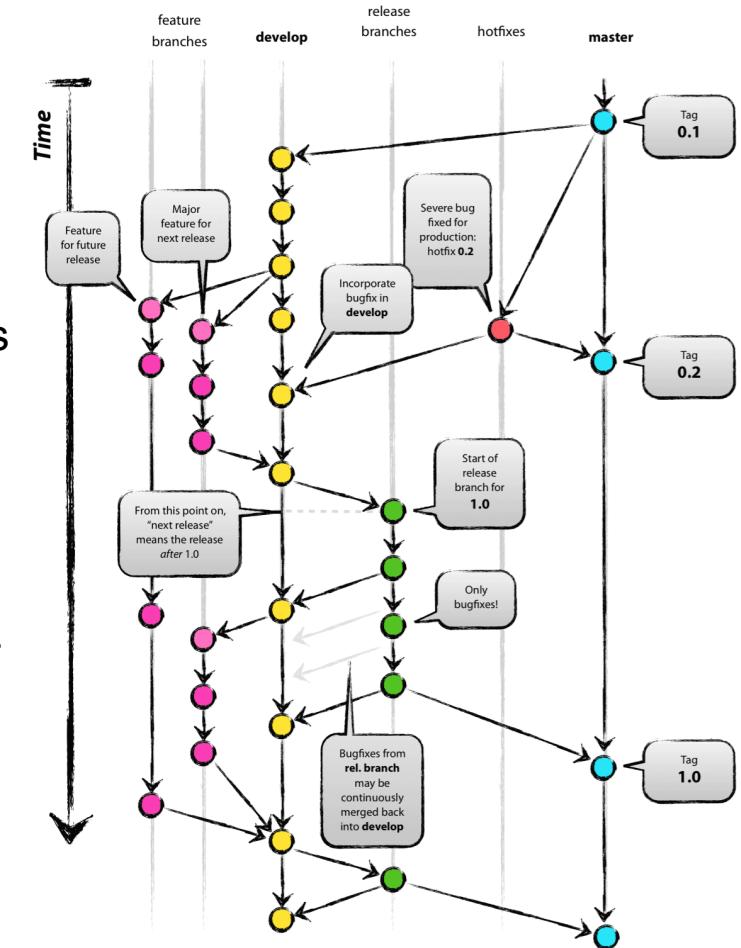
- Git Branching Model (as per nvie.com)
- Travis-CI (see travis-ci.org)

Branching model:

- Development happens on the `develop` branch
- The `master` branch is **only used for releases**, it is always production-ready
- Releases on the `master` are tagged with `vX.Y.Z` (semver)
- Branch naming:
 - **feature/** - Feature branches are prefixed (e.g. `feature/issue333-history-browser`)
 - **hotfix/** - Hotfix branches (e.g. `feature/issue333-history-browser`)

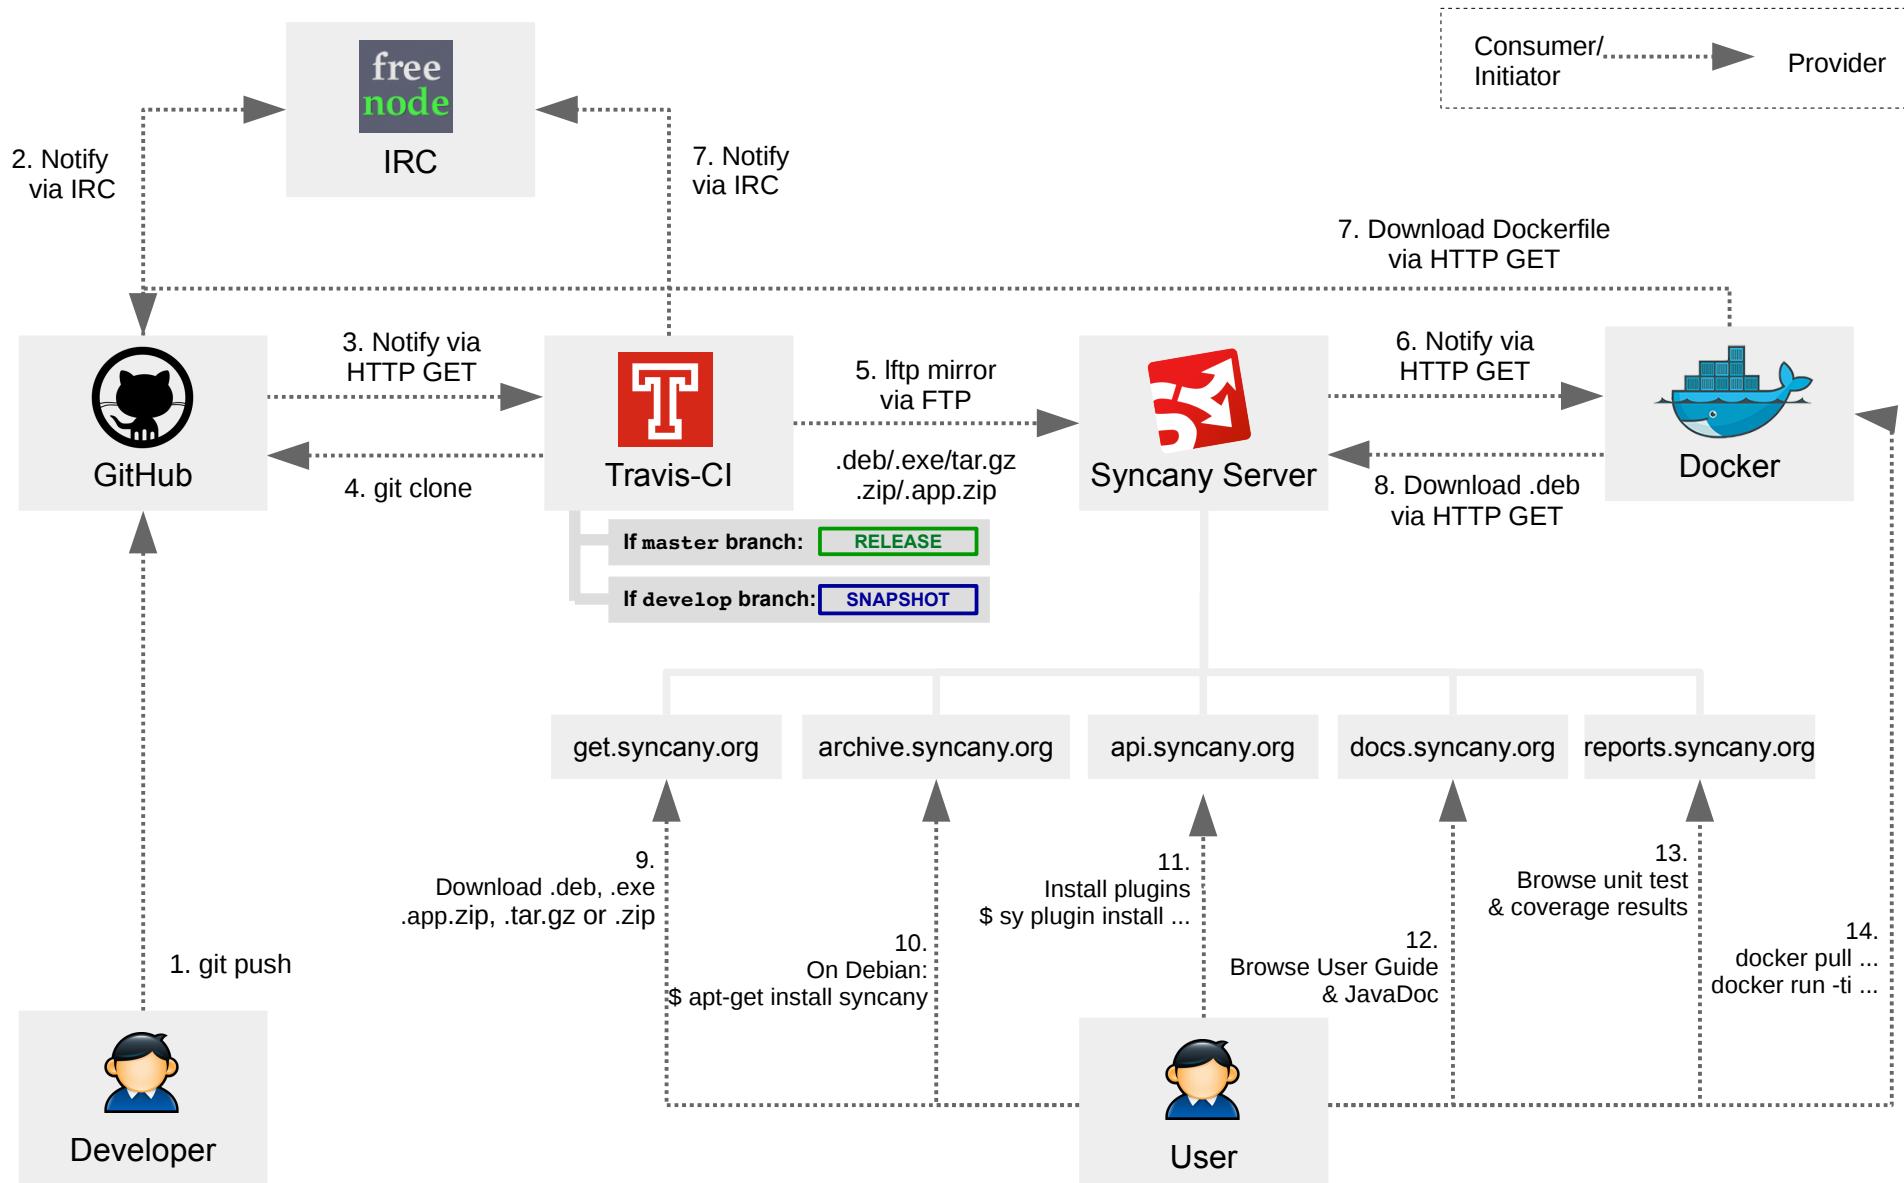
Travis CI

- Reacts to every commit
- Runs tests & builds packages, builds releases & snapshots



Continuous Integration & Packaging

> Things that happen when a developer runs “git push”



Continuous Integration & Packaging

> Debian/Windows/OSX packaging



All of the following steps are done with **every commit to the `develop` or `master` branch**. The result is either a snapshot or a new release.

Debian packaging & APT archive

- Travis: Automatic `.deb`-file packaging using `debuild`, and FTP upload to Syncany server
- Syncany server: Add to release/snapshot APT archive via `reprepro`
- Users: Download/update Syncany & plugins via `apt-get install`

Windows packaging

- Travis: Automatic creation of **Inno Setup** Windows installer (`.exe`), and FTP upload to Syncany server. *Speciality: The Inno Setup compiler is run inside Wine on Linux.*
- Syncany server: Place in release/snapshot download folder; Users: Download via website

OSX packaging

- Travis: Automatic creation of `.app.zip`-file
- Syncany server: Place in release/snapshot download folder; Users: Download via website

Other packages:

- Arch Linux PKGBUILD
- OSX homebrew recipe
- Docker image
- `.zip` and `.tar.gz` archives

Links & Resources



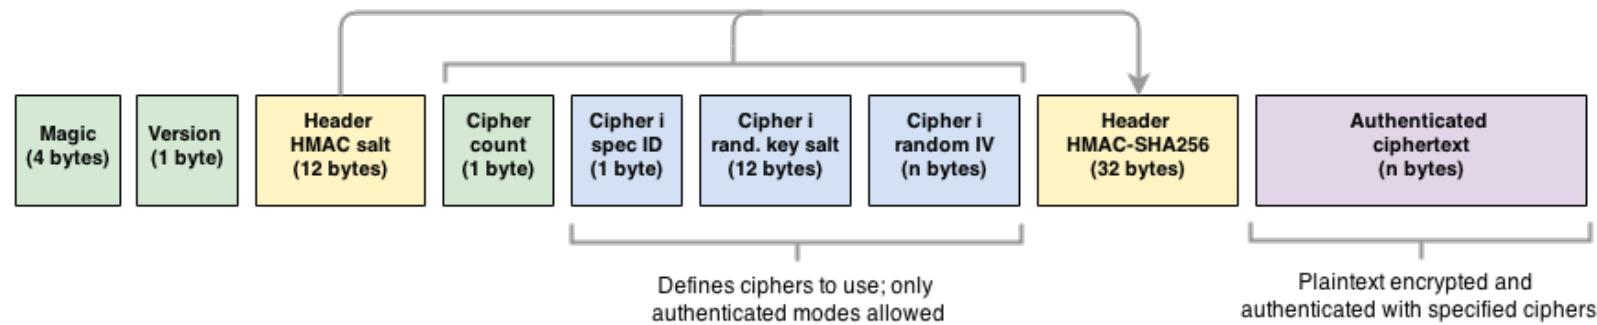
- Website - <https://www.syncany.org/>
- User Guide - <https://www.syncany.org/r/guide>
- Code & Issues - <https://github.com/syncany/syncany>
- Wiki - <https://github.com/syncany/syncany/wiki>
- IRC - #syncany on Freenode - <https://webchat.freenode.net/?channels=syncany>
(active, 20-30 people hang around here; 5-6 actually talk, my nick is "binwiederhier")
- Videos on YouTube - <https://www.youtube.com/channel/UCzegH3dpTK5HHQx6jJ5yhdQ>
- Videos on Asciinema - <https://asciinema.org/~binwiederhier>

Backup



Cryptography Concept

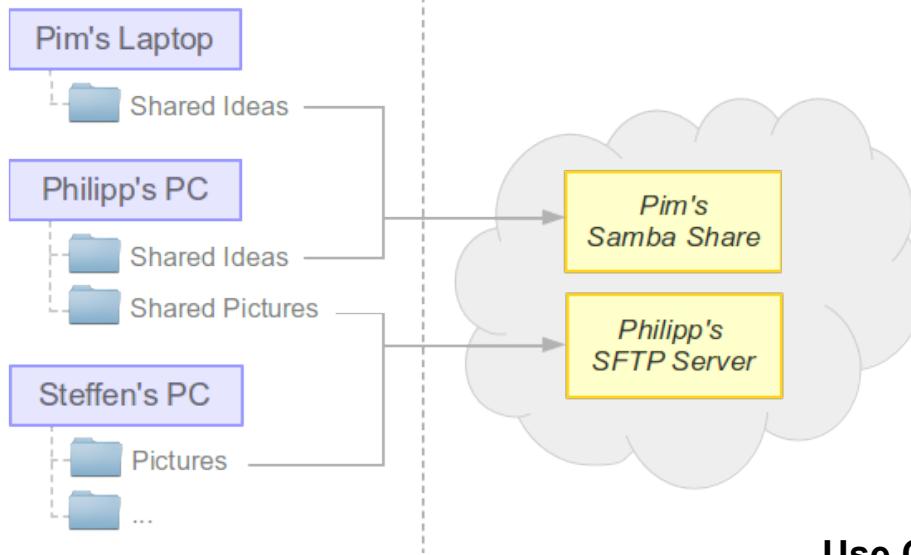
> Generic File Structure



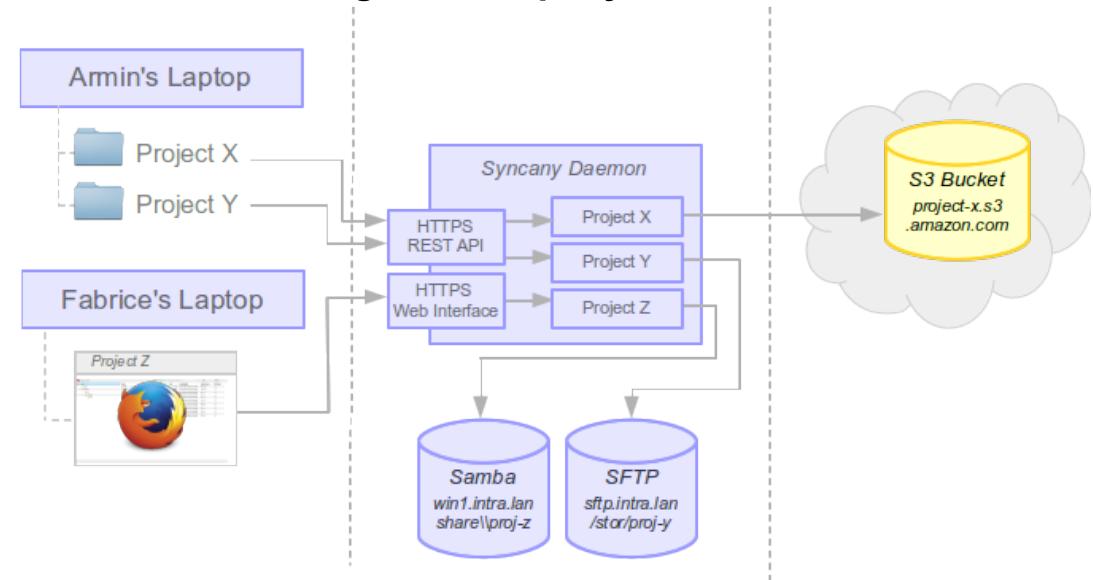


Use Cases

Use Case 1: Friends sharing files

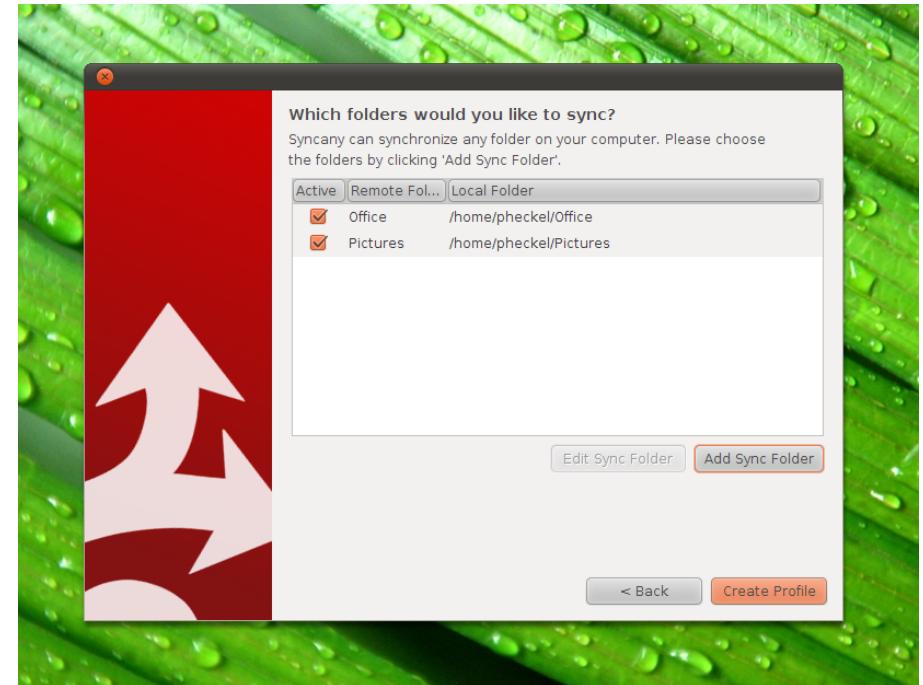
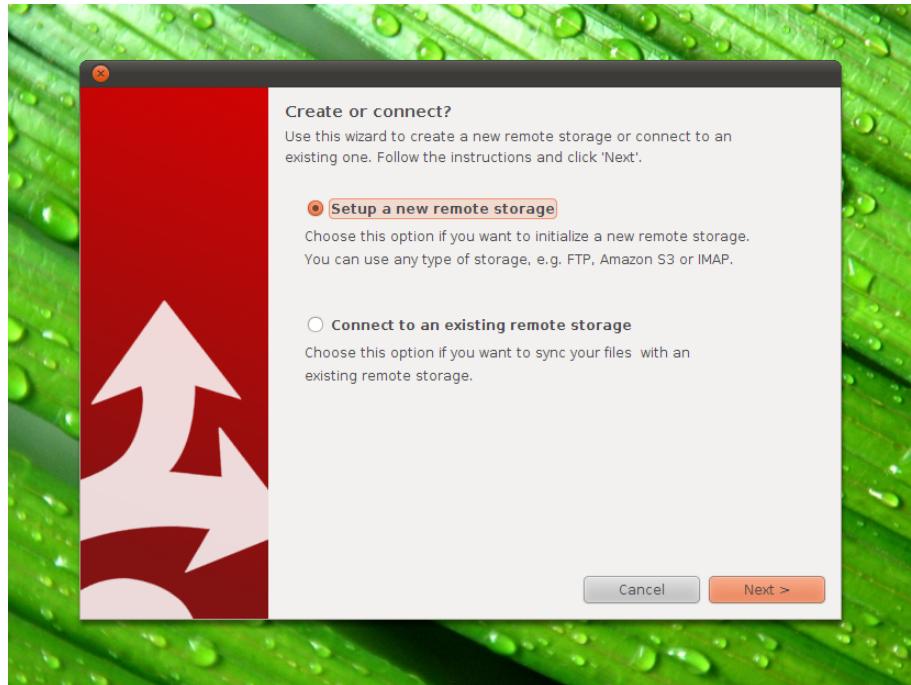


Use Case 2: Sharing in a company



Syncany Desktop User Interface

> Wizard Style Setup



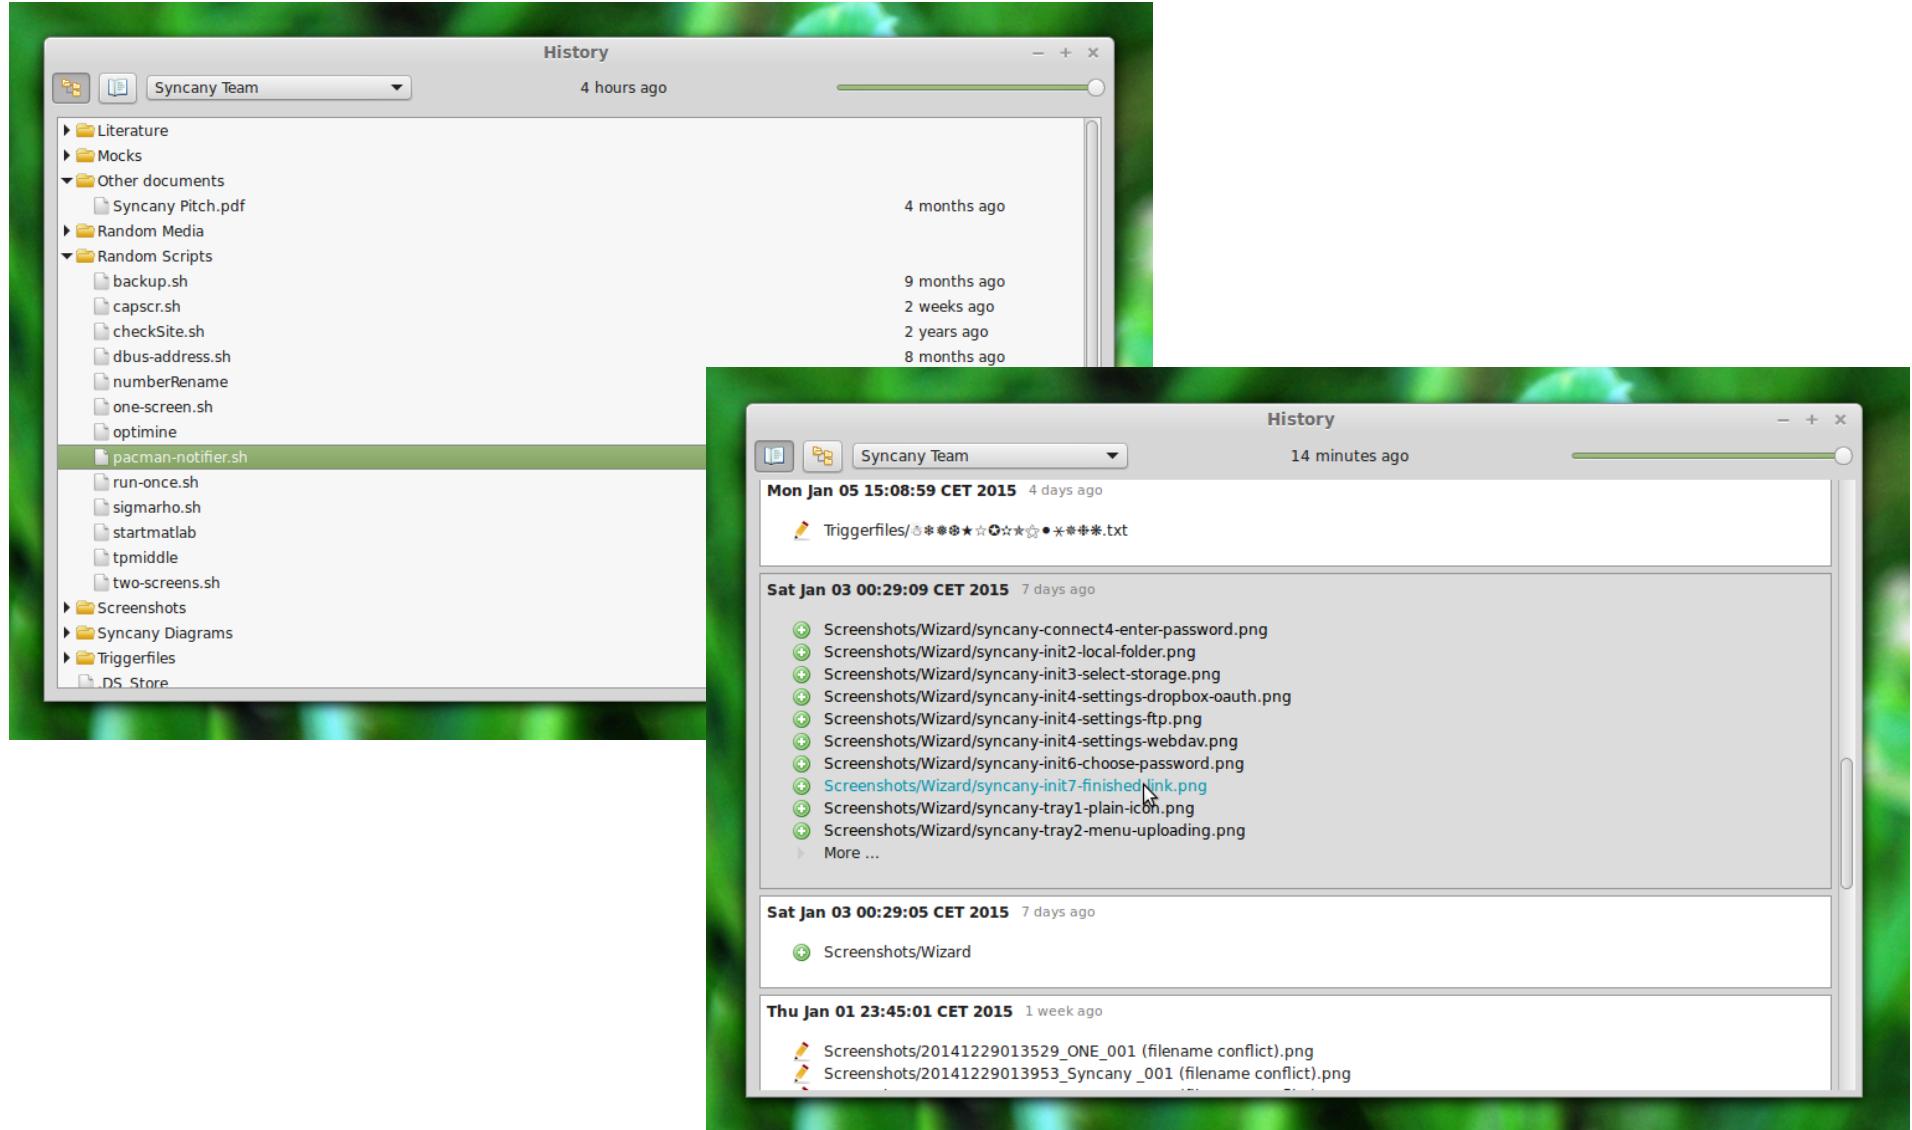
Cross-Platform Desktop Demo

https://www.youtube.com/watch?v=eHoA5_8gRBc



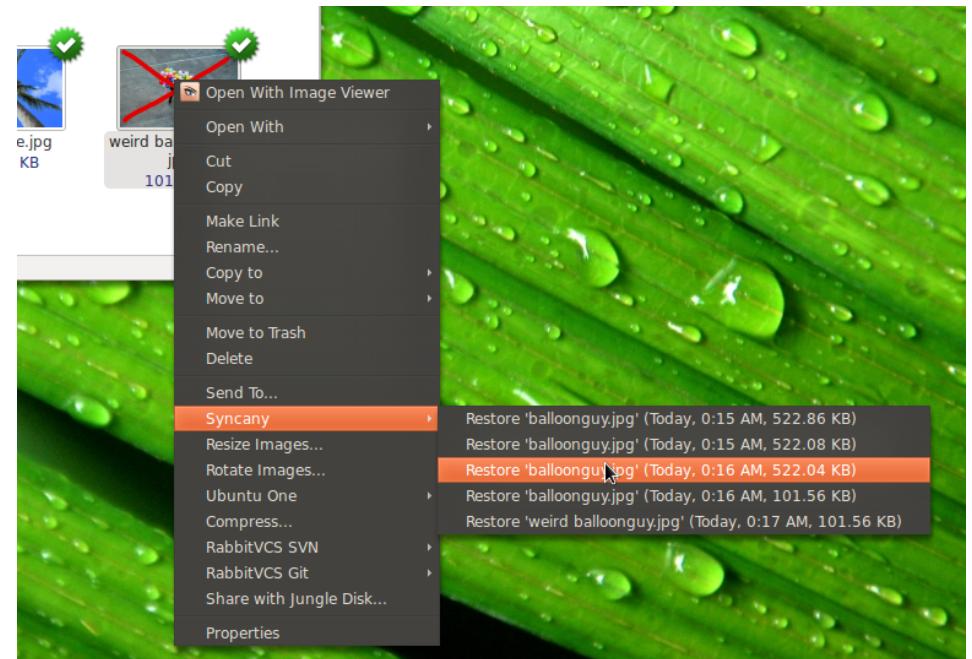
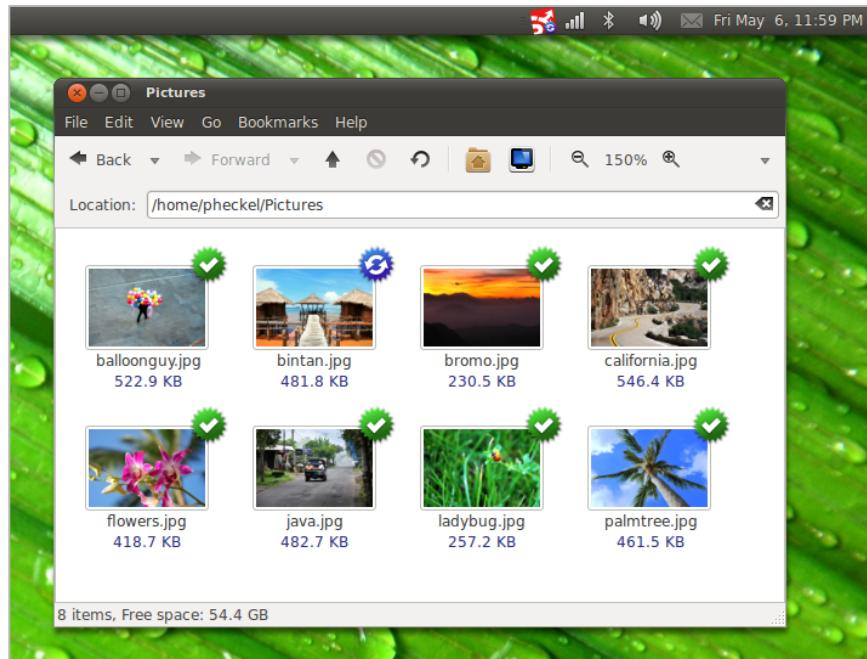
Syncany Desktop User Interface

> History browser: Browse and restore old file versions



Syncany Desktop User Interface

> File Manager Integration (2011-version PoC, looking for devs)



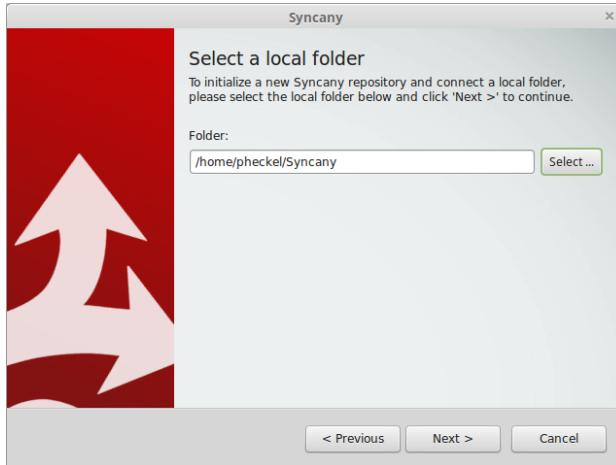
Cross-Platform Desktop Demo

https://www.youtube.com/watch?v=eHoA5_8gRBc



Using Syncany

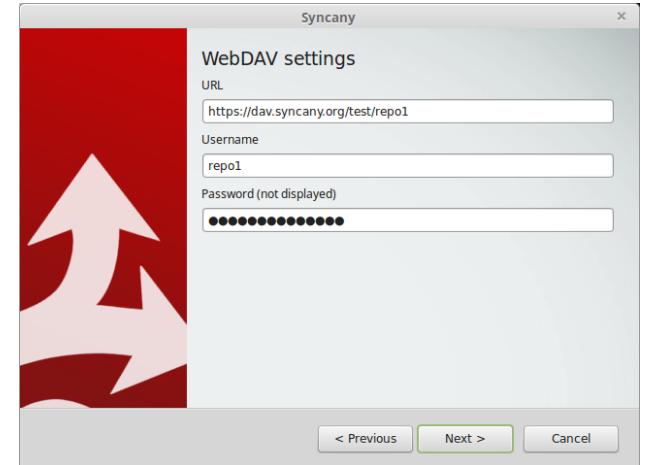
> Example: Syncing a local folder to a new WebDAV-based repository



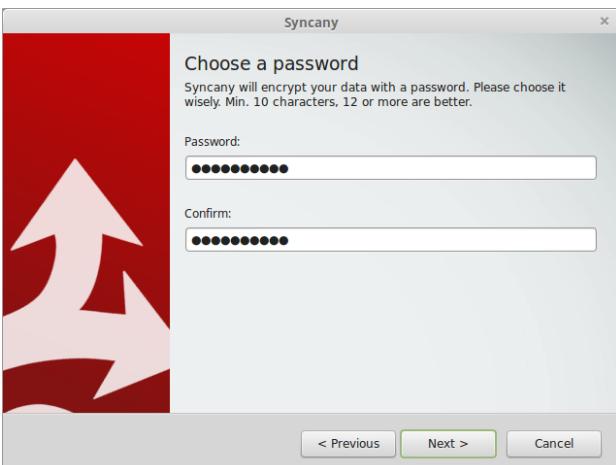
1. Choose the folder you want to sync



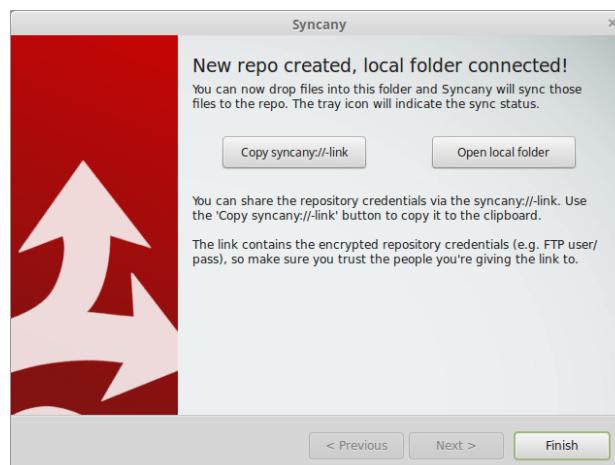
2. Choose a storage backend



3. Enter storage backend credentials



4. Choose password (used for encrypt.)



5. Done. Share a syncany://-link if you want.

Syncany Web Interface

> Web-based File Manager (proof-of-concept; looking for devs)



Syncany		Name	Version	Size	Last Modified	Posix Perms	DOS Attrs
aaa		logs	1	4.0 kB	2014-08-24 21:08:01.0 CEST	rwxr-xr-x	--a-
plugins		plugins	1	4.0 kB	2014-08-22 16:05:29.0 CEST	rwxr-xr-x	--a-
lib		Canada.jpg	1	63.3 kB	2014-08-24 02:05:48.0 CEST	rwxr--r--	--a-
userdata		daemon.ctrl	1	0 B	2014-08-24 21:08:03.0 CEST	rwxr--r--	--a-
	sftp	daemon.pid	1	5 B	2014-08-24 21:08:02.0 CEST	rwxr--r--	--a-
		daemon.xml	1	1.7 kB	2014-08-24 21:07:54.0 CEST	rwxr--r--	--a-
		daemon.xml_	1	1.4 kB	2014-08-12 23:12:44.0 CEST	rwxr--r--	--a-
		Germany.jpg	1	63.3 kB	2014-08-24 01:30:21.0 CEST	rwxr--r--	--a-
		keystore.jks	1	2.1 kB	2014-08-12 23:47:27.0 CEST	rwxr--r--	--a-
		The Netherlands.jpg	1	29.6 kB	2014-08-24 01:18:10.0 CEST	rwxr--r--	--a-
		truststore.jks	1	3.6 kB	2014-08-12 23:47:27.0 CEST	rwxr--r--	--a-
		userconfig.xml	1	414 B	2014-08-18 20:55:53.0 CEST	rwxr--r--	--a-



Web Interface Demo

<https://www.youtube.com/watch?v=2NPVffywZVs>